

Deck guide — Second Opinion

For whoever is building the slides.

Don't build a deck that explains the product. Build a deck that frames the demo, then justifies it.

Assume a **5-minute slot**: about 40 seconds before the demo, 2.5 minutes of demo, 1.5 minutes after. That's **six slides, not fifteen**. If your slot is 3 minutes, cut slides 4 and 5 and keep the rest.

THE RULES

- **Fifteen words per slide, maximum.** If a judge is reading, they aren't listening.
- **One idea per slide.** If a slide needs the word "and", split it.
- No paragraphs. No feature lists. No tech-logo soup.
- No team bios, no roadmap, no market-size slide. Nobody scores those.
- Big type — it will be projected badly in a bright room.
- Dark background matching the product, so the demo doesn't feel like a separate thing.

1 Title

10 sec

ON SCREEN

Second Opinion

Not what it is. What you're about to miss.

[team name]

YOU SAY

"We built a diagnostic co-pilot for a two-minute consultation. It never tells the doctor what the diagnosis is. It tells them what they **haven't ruled out**."

2 The problem

30 sec

ON SCREEN — ONE NUMBER, ENORMOUS

2 minutes

the average Indian primary-care consultation

Irving et al., BMJ Open 2017

YOU SAY

"That's the average consultation in India. Two minutes, in Marathi or Bhojpuri or Tamil. The doctor writes eight words of English shorthand. The patient's entire medical history is a bag of paper prescriptions nobody opens. So the symptom that would have changed the diagnosis was **said out loud and never recorded**. A better AI reading that eight-word note doesn't help — the evidence is already gone."

Why this slide exists: it makes the demo make sense, and it explains why the language layer is the product rather than decoration.

DEMO — switch to the app

2.5 min

Run the ten steps. **Do not narrate slides over it.** Land these three moments and let each one breathe:

1. **The scan.** Hold up the pill strip, photograph it, let the hormone finding appear. *"Nobody asked her that."*
2. **The reveal** (step 5). *"The differential says panic attack, 46%. The model agrees with the doctor. But look at the panel — pulmonary embolism, at under ten percent, ranked first. Because the ranking isn't probability. It's probability times what it costs to miss it. That arithmetic is printed on screen: 0.0976×100 ."*
3. **The panel empties** (step 9). *"She gets an ECG and a d-dimer, and PE leaves the list. This is a question that can be finished — which is the difference between a safety net and an alarm people learn to ignore."*

If you have time, read the Doesn't Fit line aloud: *"Panic attack does not explain hormone use — 0.0% of cases. Pulmonary embolism does, 29%."* Then: **"We didn't write that sentence. It came out of the data."**

3 How it works

30 sec

ON SCREEN — A THREE-BOX DIAGRAM, NOTHING ELSE

Sarvam captures → our code decides → Sarvam speaks
No language model assigns a number.

YOU SAY

"Speech and document AI turn the consultation and the paper into structured findings. Everything after that is counted statistics — a posterior over 49 conditions built from **1,025,602 real patient records**. The severity table is ours. The next question is information gain. Delete the AI and the reasoning still runs; you'd just have to type the codes in by hand."

Why: this answers "isn't this a ChatGPT wrapper" before anyone asks it.

4 Why the ranking is different

25 sec

ON SCREEN

risk = P × harm-if-missed

severity is log-scaled: $1 \cdot 3 \cdot 10 \cdot 30 \cdot 100$

panic attack $46\% \times 1 = 0.46$

pulmonary embolism $9.8\% \times 100 = 9.8$

YOU SAY

"Missing a heart attack isn't five times worse than missing heartburn, it's a hundred times worse. On a linear scale the likely benign thing always wins and the tool quietly agrees with whatever the doctor already thought. Log-scaled, a **10% catastrophe outranks a 46% nuisance** — which is how doctors actually triage."

5 Does it work

25 sec

ON SCREEN — TWO NUMBERS AND ONE CHART

__% top-5 recall on held-out test
__ questions to confident diagnosis vs __ at random
tuned on validate · reported on test

⚠ **Fill these in from the eval before you present.** If you don't have them, **delete this slide** — an empty metrics slide is worse than no metrics slide.

YOU SAY

"We're not claiming to be right more often than a doctor. We're claiming **coverage of the catastrophic tail**. Here's the honest number, tuned on the validation split and reported on held-out test."

6 Why it matters, and what's not done

30 sec

ON SCREEN

795,000

Americans die or are permanently disabled by diagnostic error each year
Newman-Toker et al., BMJ Quality & Safety 2023

Women are 50% more likely to be misdiagnosed after a heart attack

Leeds / MINAP, 564,412 patients, EHJ:ACC 2016

Not done yet: live streaming audio · clinician review of the severity table · calibration against Indian epidemiology

YOU SAY

"Diagnostic error is one of the largest sources of preventable harm in medicine, and it isn't evenly distributed — the case you just watched is **exactly the presentation that gets dismissed as anxiety in women**. What we haven't done: the audio is batch not streaming, no clinician has signed off our severity table yet, and our priors come from a North American dataset that would need re-weighting for India."

Why put limitations on the slide: naming your own weaknesses before a judge finds them reads as competence. Getting caught by one doesn't.

BACKUP SLIDES — AFTER THE LAST SLIDE, NEVER SHOWN UNLESS ASKED

Judges ask questions. Having the answer already on a slide is disproportionately impressive. One slide each:

- **Architecture** — the full pipeline with model names.
- **The data** — DDXPlus: 49 pathologies, 223 evidence codes, 1.03M training patients, CC-BY-4.0. Note it's synthetic, US-derived, incidence-capped.
- **Safety and regulation** — never outputs a diagnosis, never suggests treatment, always "consider" and "not yet excluded". *A pre-flight checklist, not an autopilot.*
- **Consent** — audio processed and discarded, not stored; consent is a precondition of the workflow.
- **Automation bias** — why Can't-Miss stays hidden until the doctor commits. Cite the *NEJM AI* 2026 randomised trial: physicians with 20 hours of AI training still dropped 18 points on top-choice accuracy when given flawed AI advice.
- **"Isn't this Isabel?"** — Isabel generates differentials from typed English. We rank by expected harm, compute the discriminating question, ask it aloud in the patient's language, and fit a two-minute consult.

ALSO PUT ON A SLIDE

If the rules require it: declare that the reasoning engine was built before the event, and that the hackathon work was capture, integration, evaluation and the interface. Check the rules first — but if in doubt, declare. Being upfront costs nothing; being caught costs everything.

Attribution: DDXPlus is CC-BY-4.0. One line, on the data slide or the last slide. It's a licence condition, not a courtesy.

A disclaimer in the product footer: *"Decision support only. Not a medical device. Not for clinical use."*

PRACTICAL

- **Export to PDF and keep it on the desktop.** Cloud slides fail on venue wifi.
- **Have the backup demo video one tab away**, ready if the app dies.
- **One person presents, one drives the demo.** Agree who says what and rehearse the handover — it's where teams fumble.
- **Time it out loud, five times.** Not in your head. Out loud, standing up.
- If you overrun, cut slide 5, then slide 4. **Never cut the demo.**

NUMBERS YOU'LL NEED, ALL VERIFIED

FIGURE	VALUE
Training records	1,025,602
Conditions	49
Evidence codes	223 (972 features)
Demo — differential leader	Panic attack 46.4%

Demo — panel leader	PE 9.8 = 0.0976×100
After the answer	PE 60.2%
Consultation length, India	~2 min
Bangladesh / Sweden	48 sec / 22.5 min
Diagnostic error harm, US	795,000 / year
vascular + infection + cancer	75% of it
Women post-MI misdiagnosis	~50% more likely
30-day mortality when missed	~70% higher